




tokens → concepts

VCM: Vision Concept Modeling


Adaptive Vision Token Compression via Instruction Fine-Tuning

Run Luo · Renke Shan · Longze Chen · Ziqiang Liu · Lu Wang · Min Yang · Xiaobo Xia


SIAT, Chinese Academy of Sciences · UCAS · National University of Singapore · USTC

[arXiv:2504.19627v2](https://arxiv.org/abs/2504.19627) [cs.CL]

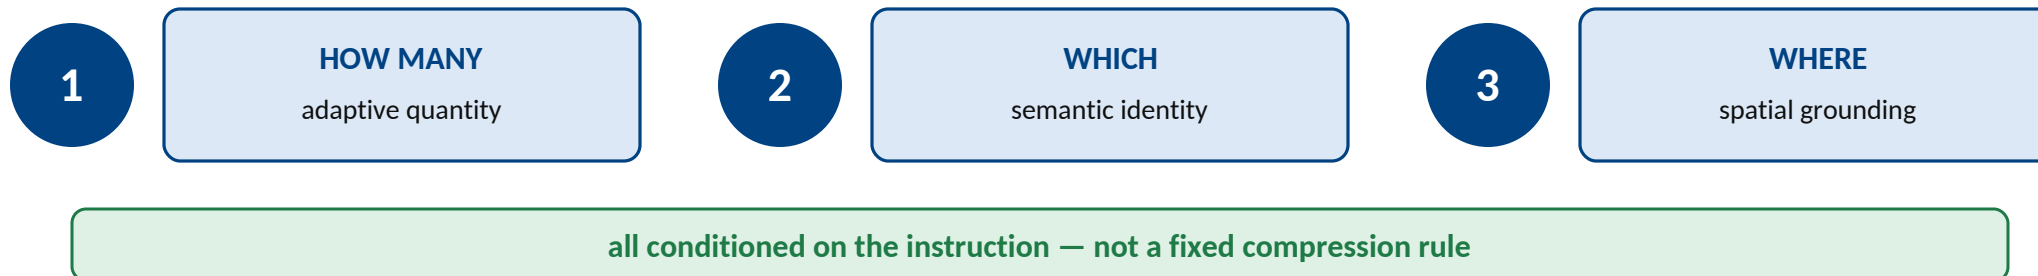
The bottleneck — LVLMs still read every image token



Compression is not yet concept modeling




A vision concept model must satisfy all three



A vision concept model must decide **how many**, **which**, and **where** — from the instruction.

VCM — one model, two training stages



Use the text prior to set a budget, then **align tokens into concepts**.

Self-supervision — the instruction tells us what vision is needed



Semantic alignment exposes the useful words

Where

is

the

person

in

yellow

?

score > mean → keyword

Masking turns one example into a free contrast

■ ■ ■ ■ ■

■ ■ ■ ■ ■

■ ■ ■ ■ ■

0 mask → 6

■ ■ ■ ■ ■

■ ■ ■ ■ ■

■ ■ ■ ■ ■

1 mask → 11

■ ■ ■ ■ ■

■ ■ ■ ■ ■

■ ■ ■ ■ ■

2 mask → 18

$$L = \text{floor}(M \cdot S \cdot [1 - \text{Norm}(N_{key})])$$

M input tokens

S information budget

Nkey keyword gap

S=1/2

288 tok

S=1/4

144 tok

S=1/6

72 tok

S=1/8


36 tok

Table 4 sweet spot

Keywords provide the signal; masking creates the contrast; **no concept labels are required.**

Optimization — learn variable-length concepts without choosing a cut

1 • Sum over every monotone alignment




2^M enumeration → dynamic programming M^2

2 • Learn the cut

$$\partial L / \partial a = p - \gamma / Z$$

prediction

posterior



3 • Merge kept runs



concept 1

concept 2


score-weighted average

parallel implementation • 100×

Unknown boundaries are marginalized during training — not guessed in advance

Forward-backward sums every alignment, then adjacent kept tokens **merge into concepts**.

Core evidence — less computation, no performance trade-off



At 144 tokens VCM beats the 576-token baseline; at $S = 1/4$ it uses **85% fewer FLOPs**.

What VCM changes — efficiency, representation, and scope

HIGH RES

2,880 → 160 tok

70.1 avg

VIDEO

2,048 → 136 tok

52.5 avg

QWEN2-VL

1,326 → 576 tok

69.6 avg

DENSE

COCO · ADE

all improve

Three contributions

- 1 Definition** quantity · identity · location
- 2 Learning** text prior + forward-backward
- 3 Outcome** efficiency + dense perception

What remains unresolved

Keyword bias

selected words may not be the true concepts

Coarse mapping

min-max length normalization is simplified

The transferable idea: let the task decide the visual granularity

Model the concepts the instruction **needs** — not every token the image **has**.