




tokens → concepts

VCM: Vision Concept Modeling


Adaptive Vision Token Compression via Instruction Fine-Tuning

Run Luo · Renke Shan · Longze Chen · Ziqiang Liu · Lu Wang · Min Yang · Xiaobo Xia

SIAT, Chinese Academy of Sciences · UCAS · National University of Singapore · USTC


[arXiv:2504.19627v2](https://arxiv.org/abs/2504.19627v2) [cs.CL]

LVLMs are the interface to the real world



Powerful, general, deployed — and still reading every image **token by token**.

Humans answer with concepts, LVLMs pay for tokens



The question needs **a few concepts**; the model still pays for **all 576 tokens**.

Why that redundancy is expensive

One LVLM forward pass, sequence length n



Attention cost is quadratic in n


$$\text{FLOPs} = T \cdot (4nd^2 + \boxed{2n^2d} + 2ndm)$$

$\text{grows with } n^2$

Cutting the vision tokens is the only lever that touches the dominant term of the sequence.

n scaled by $1/8 \rightarrow R \approx 3/25 \rightarrow 85\%$ fewer FLOPs

... and the input keeps growing




Vision tokens dominate the sequence, and attention is quadratic in its length.

Higher resolution and longer video make the **dominant** term grow **fastest**.

The existing recipe: prune, or merge

(a) Prune — threshold the attention matrix

attention matrix



top-k
→



kept tokens keep their index,
but carry no concept boundary

- no semantic concept, only surviving patches
- the number kept is set by a fixed ratio

(b) Merge — fixed-length trainable queries



vision tokens



always the same length,
original order discarded



fixed-length learnable queries


- relative order and position are lost
- length cannot follow the instruction

Both operate on tokens. Neither returns a concept that a downstream task can point at in the image.

Both make the sequence **shorter**; neither makes it **conceptual**.

What is still missing

No spatial grounding



compressed vectors do not say
where in the image they came from

Order and position lost




order no longer recoverable

→ region-level tasks break

merged queries drop the relative
layout the answer may depend on

Length not controllable




a fixed ratio cannot follow what
the instruction actually asks for

Consequence: applicable to visual question answering, and little else.

Detection, segmentation and region-level tasks all need concepts that stay attached to the image.

Shorter is not the same as conceptual — so none of them is a vision concept model.

Definition: the vision concept model




A model that decides **how many** concepts, **which** ones, and **where** — from the instruction.

Two challenges stand in the way

Challenge 1 • no concept-level labels

abundant
VQA instruction data



no concept annotation

concept-level labels
would be exact — but

labor-intensive
time-consuming
does not scale

so nobody has them


Needed: a signal already inside the data.

Challenge 2 • the length is not fixed

“Where is the person in yellow?”

“What is the orange number of the bus?”

“Can you describe this image in detail?”



3

5


10

concepts actually needed

searching every cut point: $O(2^M)$ $M = 576$

We need supervision **without labels**, and optimization **without a fixed length**.


The text prior already knows the answer



Probe: 5K LLaVA instances • GPT-4o extracts image-related keywords • VisionZip at 24 lengths • GPT-4o judges the minimum sufficient length.

The instruction-response **keyword gap** is a supervision signal that costs nothing.


VCM at a glance



Select keywords, estimate a length, then **align by dynamic programming**.

Figure 3(a) of the paper.


Step 1 — keywords from semantic alignment



Semantic alignment loss

$$L_{SA} = -p(G^T_{LLM}) \log p(G^V) - p(G^T_{LLM}) \log p(G^T)$$

pulls text, vision and LLM features into one space




$$K = \text{Softmax}(\text{MHSA}([H^I; H^R]) \cdot G^V), \quad \text{keep } k > E(K)$$

→ “person” and “yellow” survive; the rest is function words

Keywords fall out of **vision-language alignment**, not out of human annotation.

Step 2 — implicit contrastive sampling



more masking → more vision information required


Why this is safe

Masking never shortens the requirement, so the objective cannot learn to discard information the answer depends on.

$$r \sim \text{Uniform}(0, 1)$$

$$\epsilon(r) = a + (b-a) \cdot [1 + \tanh(k(2r-1))] / 2$$

$$a = 0.2, b = 1.2, k = 5$$



Masking keywords **lengthens** what is needed — a free contrast, and a **knob** at inference.

Step 3 — turn the prior into a target length

$$L = \text{floor}(M \cdot S \cdot [1 - \text{Norm}(N_{\text{key}})])$$


M input vision tokens **S** information domain **N_{key}** keyword gap

The information domain sets the budget

S	max tokens	Avg. (%)
1/2	288	58.9
1/4	144	59.3
1/6	72	56.4
1/8	36	57.2

Avg. over 4 VQA benchmarks (Table 4).


The keyword gap moves the target length



Exact evaluation of the formula at M = 576; S = 1/4; N_{key} is clipped to [-35, +10].

The text prior becomes an explicit **target length** — still no concept labels.

Step 4 — forward-backward over all alignments



exhaustive search
 $O(2^M) \rightarrow O(M^2)$
by dynamic programming

$$\alpha(t, l) = p(z_l | y_t) \cdot [\alpha(t-1, l) + \alpha(t-1, l-1)]$$

$$\beta(t, l) = [\beta(t+1, l) + \beta(t+1, l+1)] \cdot p(z_l | y_t)$$

Every monotone path through the lattice is one way to cut the sequence into L concepts; the forward and backward recursions score all of them at once.

Where to cut is unknown, so we sum over **every** alignment: $2^M \rightarrow M^2$.

The gradient is a posterior residual

$\gamma = \alpha\beta / p(Z|Y)$ — the posterior over alignments

	l=1	l=2	l=3	l=4	l=5	row sum
t1	1.818	0.196	0.000	0.000	0.000	2.014
t2	1.163	0.819	0.032	0.000	0.000	2.014
t3	0.248	1.677	0.074	0.015	0.000	2.014
t4	0.075	1.447	0.464	0.027	0.002	2.014
t5	0.052	0.458	1.449	0.042	0.014	2.014
t6	0.047	0.051	1.723	0.167	0.027	2.014
t7	0.005	0.088	0.177	1.708	0.036	2.014
t8	0.002	0.046	0.133	0.943	0.890	2.014

The row sum is constant in t — every column redistributes the same total probability mass.

From posterior to gradient

$$p(Z^V|Y^V) = \alpha(M,2L) + \alpha(M,2L+1) = \beta(1,1) + \beta(1,2)$$



$$L_{VCM} = -\log p(Z^V|Y^V)$$



$$\partial L_{VCM} / \partial y_t(z_l) = p(z_l|y_t) - \gamma(t,l)$$


prediction minus posterior — the same shape as CTC.

keep token t when $p(\bullet | y) > p(\emptyset | y)$

Worked example from Table 6: $M = 8$ input tokens, target length $L = 2$. Row sums of γ stay constant across t .

One clean rule — prediction minus posterior, exactly CTC's gradient shape.

Step 5 — merge the kept runs into concepts



Implementation

Variable-length segments are merged with cumulative sums, not a double loop over batch and sequence.

100×
faster

The objective actually optimized

$$L^I = L_{NTP} + \alpha \cdot L_{SA}$$



pre-training

$$L^{II} = L_{NTP} + \varepsilon(r) \cdot L_{VCM}$$

instruction fine-tuning

Adjacent kept tokens merge by **score-weighted average** — and the op is **100× faster** than a double loop.

Evidence — better than the model it compresses




Average over 11 image-based VQA benchmarks (Table 1).

144 tokens instead of 576, and a higher average than LLaVA-1.5 itself.

Table 1 of the paper. #Vision Tokens = tokens fed to the LLM backbone.

Evidence — the cost actually drops




Average over the 8 benchmarks reported in Table 10; FLOPs and latency measured in the same setting.


85% fewer FLOPs in theory, 1.8× lower latency in practice, no average loss.

Table 10 and Appendix D of the paper. Analytic ratio $R \approx 3/25$ at $S = 1/4$.

Evidence — the vision encoder itself improves




Why it transfers



CLIP ViT-L/14 with and without VCM fine-tuning, dropped into F-VLM (detection) and Cat-Seg (segmentation).

Not only compression — concept modeling improves dense perception.


Evidence — the recipe transfers




High resolution (Table 7) • video (Table 8) • other architectures and larger LLMs (Table 9).

Resolution, modality, architecture, scale — **the same recipe holds.**

Ablation — what each piece buys



Avg. (%) — vertical axis starts at 54.0



Avg. (%) — vertical axis starts at 54.0

Average over 4 image-based VQA benchmarks, batch size 128, 500 training steps (Table 4).

Every component contributes, and $S = 1/4$ is the accuracy–cost sweet spot.

Conclusion

Contributions

- 1 Define the vision concept model**
quantity, identity and location, all conditioned on the instruction
- 2 A self-supervised framework**
keyword selection plus a forward-backward algorithm; no concept labels
- 3 Efficiency and capability together**
85% fewer FLOPs, and a stronger dense-perception encoder

Limitations the paper states

Keyword selection is adaptive

selected keywords may not be the true ones — a source of bias

Length estimate is coarse

min-max normalization is a simplification of the mapping

Next: finer keyword selection with open-source LLMs.

Model the concepts the instruction needs — not every token the image has.

Model the concepts the instruction **needs** — not every token the image **has**.